

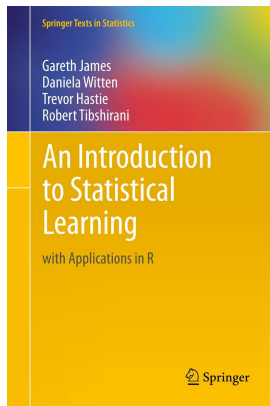
Model selection

Timothy Daley

November 29, 2016

Book change for this week

This week we will be using the book *An Introduction to Statistical Learning with Applications in R*, available for free at [<http://www-bcf.usc.edu/~gareth/ISL/>]. The material for this week is covered in chapter 6.



Multiple regression: regress everything!

- Suppose we have a bunch of data and we just run a regression on everything.

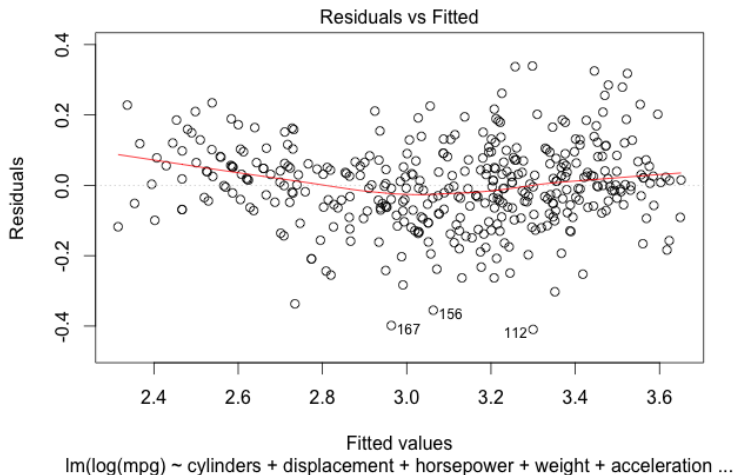
```
log_mpg_vs_all.lm = lm(log(mpg) ~ cylinders + displacement  
                        + horsepower + weight  
                        + acceleration + year  
                        + origin, data=Auto)
```

Multiple regression: regress everything!

Table 1: Fitting linear model: $\log(\text{mpg}) \sim \text{cylinders} + \text{displacement} + \text{horsepower} + \text{weight} + \text{acceleration} + \text{year} + \text{origin}$

	Estimate	Std. Error	t value	Pr(> t)
cylinders	-0.02795	0.01157	-2.415	0.01619
displacement	0.0006362	0.000269	2.365	0.01852
horsepower	-0.001475	0.0004935	-2.989	0.002984
weight	-0.0002551	2.334e-05	-10.93	2.12e-24
acceleration	-0.001348	0.003538	-0.381	0.7034
year	0.02958	0.001824	16.21	2.13e-45
origin	0.04071	0.009955	4.089	5.276e-05
(Intercept)	1.751	0.1662	10.53	5.845e-23
R^2	0.87951			

Multiple regression: regress everything!



How to choose which variables are important?

- Method 1: Successively remove variables that are not important.

	Est	imate Std	. Error t v	alue Pr(> t)
cylinders	-0.02795	0.01157	-2.415	0.01619*
displacement	0.0006362	0.000269	2.365	0.01852*
horsepower	-0.001475	0.0004935	-2.989	0.002984*
weight	-0.002551	2.334e-05	-10.93	2e-24*
acceleration	-0.001348	0.003538	-0.381	0.7034
year	0.02958	0.001824	16.21	3e-45*
origin	0.04071	0.009955	4.089	5.27e-05*
(Intercept)	1.751	0.1662	10.53	5.845e-23*
R^2	0.87951			

Removing non-significant variables

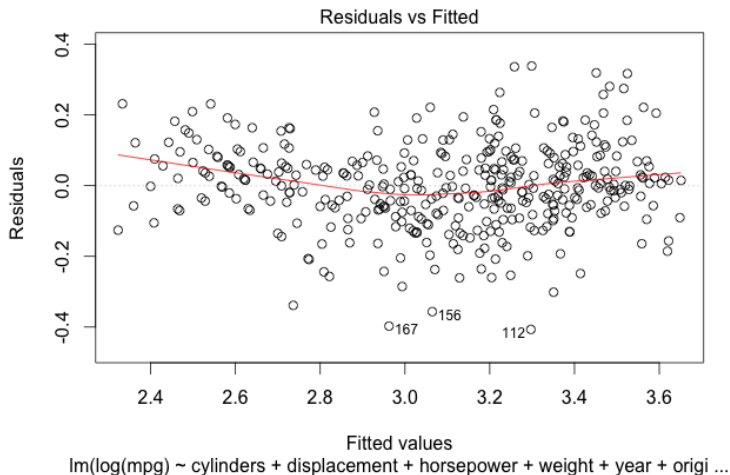
```
log_mpg_vs_sig.lm.1 = lm(log(mpg) ~ cylinders + displacemen  
+ horsepower + weight + year  
+ origin, data = Auto)
```

Removing non-significant variables

Table 3: Fitting linear model: $\log(\text{mpg}) \sim \text{cylinders} + \text{displacement} + \text{horsepower} + \text{weight} + \text{year} + \text{origin}$

	Estimate	Std. Error	t value	Pr(> t)
cylinders	-0.02772	0.01154	-2.402	0.01679*
displacement	0.0006466	0.0002673	2.419	0.01601*
horsepower	-0.001359	0.0003876	-3.505	0.0005099*
weight	-0.0002594	2.044e-05	-12.69	4.433e-31*
year	0.02963	0.001817	16.31	7.782e-46*
origin	0.04067	0.009944	4.09	5.247e-05*
(Intercept)	1.723	0.1493	11.54	1.174e-26*
R^2	0.87946			

Removing non-significant variables



How to choose which variables are important?

- ▶ Method 2:
 - ▶ Start with nothing (intercept only)
 - ▶ Test adding variables one at a time
 - ▶ Add most significant
 - ▶ Repeat until no variables are significant

Adding significant variables

- Individual regression estimates:

	Estimate	Std. Error	t value	Pr(>
cylinders	-0.1648	0.005678	-29.03	1.675e-99
displacement	-0.00277	8.569e-05	-32.37	1.478e-112
horsepower	-0.00733	0.0002494	-29.41	5.393e-101
weight	-0.00035	9.79e-06	-35.81	2.392e-125*
acceleration	0.05516	0.005581	9.884	1.045e-20
year	0.05329	0.003817	13.96	3.28e-36
origin	0.2366	0.0177	13.37	8.27e-34

Adding significant variables

- Add weight, now the other added regression coefficients are:

	Estimate	Std. Error	t value	Pr(>
cylinders	-0.04193	2.182e-05	-12.6	8.793e-31
displacement	-0.00092	0.000216	-4.262	2.549e-05
horsepower	-0.00256	0.00041	-6.231	1.205e-09
acceleration	0.01232	0.003261	3.777	0.0001835
year	0.03129	0.00177	17.68	8.933e-52*
origin	0.03096	0.01265	2.448	0.01481

Adding significant variables

- Add weight and year, now the other added regression coefficients are:

	Estimate	Std. Error	t value	Pr(>
cylinders	-0.01872	0.00831	-2.253	0.02483
displacement	-0.00024	0.000169	-1.429	0.1537
horsepower	-0.00087	0.000335	-2.612	0.009355
acceleration	0.0043	0.00251	1.714	0.08735
origin	0.03081	0.009372	3.288	0.001103*

Adding significant variables

- Add weight, year, and origin, now the other added regression coefficients are:

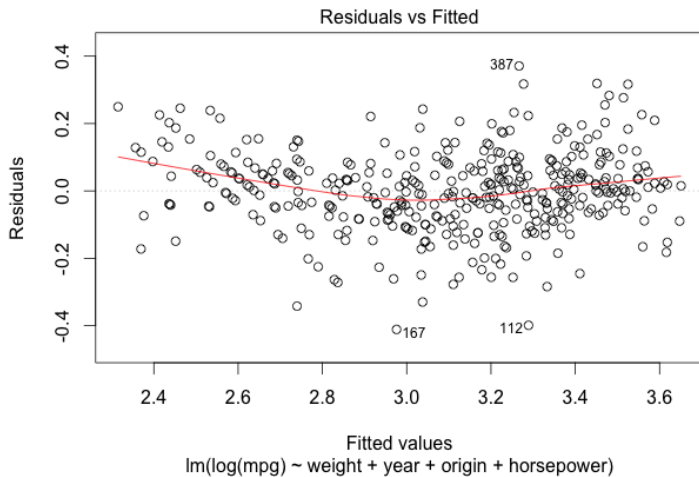
	Estimate	Std. Error	t value	Pr(>
cylinders	-0.01559	0.008287	-1.882	0.06064
displacement	-0.00012	0.000173	-0.6686	0.5042
horsepower	-0.00104	0.000332	-3.116	0.001968*
acceleration	0.00466	0.002479	1.881	0.06077

Adding significant variables

- Final model: weight, year, origin, and horsepower

	Estimate	Std. Error	t value	Pr(>
weight	-0.00025	1.58e-05	-15.89	3.804e-44
year	0.0295	0.001821	16.2	1.941e-45
origin	0.03463	0.009349	3.704	0.0002434
horsepower	-0.001034	0.0003323	-3.116	0.001968
(Intercept)	1.658	0.148	11.2	1.989e-25

Adding significant variables



Adding vs subtracting variables

- ▶ Subtracting variables:
 - ▶ Final model:
 - ▶ $\log(\text{mpg}) \sim \text{cylinders} + \text{displacement} + \text{horsepower} + \text{weight} + \text{year} + \text{origin}$
- ▶ Adding variables:
 - ▶ Final model:
 - ▶ $\log(\text{mpg}) \sim \text{weight} + \text{year} + \text{origin} + \text{horsepower}$
- ▶ Which model is better?

Adjusted R^2

- ▶ R^2 always increases with the number of variables
- ▶ Penalize R^2 for the number of variables
- ▶ adjusted $R^2 = R^2 - (1 - R^2) \frac{p}{n-p-1}$
 - ▶ n is the number of observations (sample size)
 - ▶ p is the number of variables in regression

Mallow's C_p

- ▶ $C_p = \frac{1}{n}(RSS + 2p\hat{\sigma}^2)$
 - ▶ RSS = residual sum of squares = $\sum_{i=1}^n (y_i - \hat{y}_i)^2$
 - ▶ $\hat{\sigma}^2$ is the estimated variance of the residuals
- ▶ Smaller C_p is better.

AIC: Akaike Information Criterion

- ▶ In general: $AIC = 2p - 2 \log L$
- ▶ For linear regression: $AIC = 2p + n \log RSS$
- ▶ With more predictors, RSS gets smaller while p gets larger
- ▶ Smaller AIC is preferable
- ▶ C_p is a special case of AIC.

BIC: Bayesian Information Criterion

- ▶ In general: $\text{BIC} = p \log n - 2 \log L$
- ▶ For linear regression: $\text{BIC} = p \log n + n \log(RSS/n)$
- ▶ Like AIC but larger penalty on number of parameters
 - ▶ Tends to pick smaller models than AIC

Regularization methods

- ▶ In linear models the goal is to minimize the residual sum of squares

$$RSS = \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2$$

- ▶ What if we add a penalty term for the β_j 's?
 - ▶ Don't include β_0 since we could just set $\beta_0 = \bar{y}$ and de-mean the y 's

Ridge regression

- ▶ Ridge regression adds a penalty proportional to β_j^2 :

$$\sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2$$

- ▶ This “shrinks” the β_j ’s towards 0
 - ▶ But not all the way.

LASSO (least absolute shrinkage and selection operator)

- ▶ Lasso adds a penalty proportional to $|\beta_j|$:

$$\sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\beta_j|$$

- ▶ This “shrinks” the β_j ’s to be 0

How to choose λ ?

- ▶ $\lambda = 0$ is equivalent to plain regression
- ▶ $\lambda = \infty$ will make all $\beta_j = 0$.
- ▶ Want $0 < \lambda < \infty$, but what is optimal?