

Cross Validation

Timothy Daley

November 30, 2016

How to test models on the data?

- ▶ Sometimes we can build a model, then get more data to test the model.
- ▶ Most times we can't. We have the data we have.
- ▶ What to do?

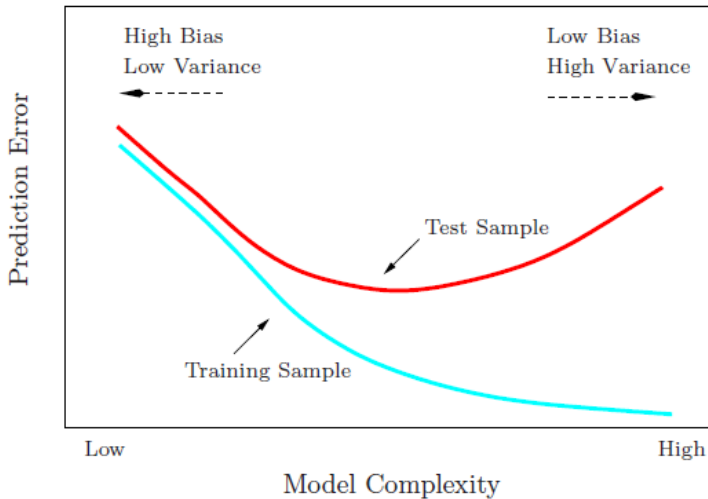
Hold-out

- ▶ Take some of the data and sequester it for testing
- ▶ Building the model ignores this data.
 - ▶ Can introduce bias if the test data is poorly selected.
- ▶ How to use the full data for training and testing?

Cross validation

- ▶ Divide the data into two non-overlapping parts:
 - ▶ Training set to build the model
 - ▶ Test set to test or validate the model
 - ▶ The model will naturally fit the training set better than the test set.
 - ▶ Use the performance test set to choose the best model
- ▶ Training error is $RSS/n = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$ for y_i in training set
- ▶ Test error is error in test set $\frac{1}{m} \sum_{i=1}^m (y_i - \hat{y}_i)^2$

Training vs Test Error performance



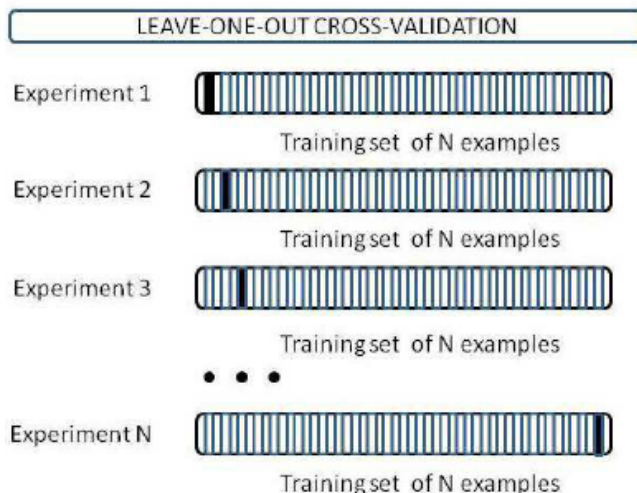
Cross validation

- ▶ In cross validation you successively divide the data into training and test sets
- ▶ Average test error over all iterations
- ▶ Choose model with lowest average test error
- ▶ Refit model using full data

Leave one out cross validation

- ▶ Leave one out cross validation:
 - ▶ For each of the n observations, take the test set to be a single observation and the training set to be the other $n - 1$ observations.
 - ▶ Average performance across the test sets.
 - ▶ n total test iterations

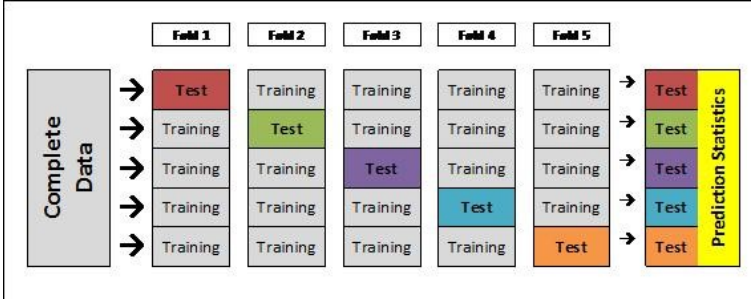
Leave one out cross validation



k -fold cross validation

- ▶ In k -fold cross validation you divide the data into k parts.
- ▶ Use each k parts as test sets, successively, and the remainder as training.
- ▶ Average performance across k test sets.

k-fold cross validation



Cross validation in practice

- ▶ Let's apply cross validation to evaluate the performance of the forward and backward stepwise regression models we built last class.
- ▶ Forwards model: $\log(\text{mpg}) \sim \text{weight} + \text{year} + \text{origin} + \text{horsepower}$
- ▶ Backwards model: $\log(\text{mpg}) \sim \text{cylinders} + \text{displacement} + \text{horsepower} + \text{weight} + \text{year} + \text{origin}$

Cross validation in practice

```
n_folds = 10
fold = sample(1:n_folds, dim(Auto)[1],
              replace = TRUE)
forward_error = 0
backward_error = 0
```

Cross validation in practice

```
for(i in 1:n_folds){  
  forward.lm = lm(log(mpg) ~ weight + year  
                  + origin + horsepower,  
                  data = Auto[-which(fold == i), ])  
  forward_error = forward_error +  
    mean((log(Auto$mpg[which(fold == i)]) -  
          predict(forward.lm,  
                  newdata = Auto[which(fold == i), ],  
                  interval = "none"))^2)  
}
```

Cross validation in practice

```
for(i in 1:n_folds){  
  backward.lm = lm(log(mpg) ~ cylinders + displacement  
                    + horsepower + weight + year + origin,  
                    data = Auto[-which(fold == i), ])  
  backward_error = backward_error +  
    mean((log(Auto$mpg[which(fold == i)]) -  
          predict(backward.lm,  
                  newdata = Auto[which(fold == i), ],  
                  interval = "none"))^2)  
}
```

Cross validation in practice

```
forward_error
```

```
## [1] 0.1469134
```

```
backward_error
```

```
## [1] 0.1456321
```

Backwards error is lower, use larger model: $\log(\text{mpg}) \sim \text{cylinders} + \text{displacement} + \text{horsepower} + \text{weight} + \text{year} + \text{origin}$

Comparison with AIC, BIC, and adjusted R^2

```
backward.lm = lm(log(mpg) ~ cylinders + displacement  
                  + horsepower + weight + year + origin,  
                  data = Auto)  
forward.lm = lm(log(mpg) ~ weight + year + origin  
                 + horsepower, data = Auto)
```


Comparison with AIC, BIC, and adjusted R^2

```
library(MASS)  
AIC(forward.lm)
```

```
## [1] -544.6394
```

```
AIC(backward.lm)
```

```
## [1] -547.6487
```

Comparison with AIC, BIC, and adjusted R^2

```
BIC(forward.lm)
```

```
## [1] -520.8118
```

```
BIC(backward.lm)
```

```
## [1] -515.8786
```

Comparison with AIC, BIC, and adjusted R^2

```
summary(forward.lm)$adj.r.squared
```

```
## [1] 0.8760216
```

```
summary(backward.lm)$adj.r.squared
```

```
## [1] 0.8775862
```