

Model Selection in practice

Timothy Daley

November 30, 2016

Dataset: Wage

- For this example I'm going to be looking at the Wage dataset in the ISLR package.

```
library(ISLR)
library(pander)
data(Wage)
library(glmnet)
library(MASS)
```

Wage

```
names(Wage)
```

```
## [1] "year"      "age"       "sex"       "maritl"  
## [6] "education" "region"    "jobclass"  "health"  
## [11] "logwage"   "wage"
```

```
levels(Wage$education)
```

```
## [1] "1. < HS Grad"      "2. HS Grad"        "3. Some C  
## [4] "4. College Grad"   "5. Advanced Degree"
```

Applying the lasso to predict wage: prepping the data

```
xfactors = model.matrix(wage ~ sex + maritl + race +  
                        education + region + jobclass +  
                        health + health_ins,  
                        data = Wage)[-1]  
x = as.matrix(data.frame(year = Wage$year, age = Wage$age,  
                        xfactors))
```

Using Cross-Validation to choose λ

- We use the function `cv.glmnet`

```
?cv.glmnet
```

```
cv.glmnet {glmnet}
```

R Documentation

Cross-validation for glmnet

Description

Does k-fold cross-validation for glmnet, produces a plot, and returns a value for `lambda`

Usage

```
cv.glmnet(x, y, weights, offset, lambda, type.measure, nfolds, foldid, grouped, keep,
parallel, ...)
```

Arguments

<code>x</code>	<code>x</code> matrix as in <code>glmnet</code> .
<code>y</code>	response <code>y</code> as in <code>glmnet</code> .
<code>weights</code>	Observation weights; defaults to 1 per observation
<code>offset</code>	Offset vector (matrix) as in <code>glmnet</code>
<code>lambda</code>	Optional user-supplied <code>lambda</code> sequence; default is <code>NULL</code> , and <code>glmnet</code> chooses its own sequence
<code>nfolds</code>	number of folds - default is 10. Although <code>nfolds</code> can be as large as the sample size (leave-one-out CV), it is not recommended for large datasets. Smallest value allowable is <code>nfolds=3</code>


Using Cross-Validation to choose λ

```
# lasso is alpha = 1  
wage.cvfit = cv.glmnet(x, Wage$wage, alpha = 1)  
wage.cvfit$lambda.min
```

```
## [1] 0.0285344
```

Using Cross-Validation to choose λ

```
plot(wage.cvfit)
```



Cross-Validated model

```
head(coef(wage.cvfit, s = "lambda.min"), 13)
```

```
## 13 x 1 sparse Matrix of class "dgCMatrix"
```

```
##                                1
```

```
## (Intercept)                -2392.1122077
```

```
## year                        1.2260075
```

```
## age                         0.2710652
```

```
## sex2..Female                .
```

```
## maritl2..Married            17.0423512
```

```
## maritl3..Widowed           1.5232654
```

```
## maritl4..Divorced           3.7478013
```

```
## maritl5..Separated          11.0884884
```

```
## race2..Black                -5.0042586
```

```
## race3..Asian                -2.6717244
```

```
## race4..Other                -5.9340531
```

```
## education2..HS.Grad         7.1210700
```

```
## education3..Some.College    17.6820562
```


Cross-Validated model

```
tail(coef(wage.cvfit, s = "lambda.min"), 13)
```

```
## 13 x 1 sparse Matrix of class "dgCMatrix"
##                                     1
## education4..College.Grad          30.592752
## education5..Advanced.Degree       53.288490
## region2..Middle.Atlantic           .
## region3..East.North.Central         .
## region4..West.North.Central         .
## region5..South.Atlantic             .
## region6..East.South.Central         .
## region7..West.South.Central         .
## region8..Mountain                  .
## region9..Pacific                   .
## jobclass2..Information              3.556992
## health2....Very.Good               6.484683
## health_ins2..No                    -17.533736
```

Why is sex not included?

```
table(Wage$sex)
```

```
##
```

```
##  1. Male 2. Female
```

```
##      3000          0
```

Oh. There's no females in the data.

Predicting MLB salary

The hitters data set in ISLR has data from the 1986 and 1987 seasons. Let's see what individual statistics most contribute to salary

```
data(Hitters)
names(Hitters)
```

```
## [1] "AtBat"      "Hits"       "HmRun"      "Runs"       "RBI"
## [6] "Walks"      "Years"      "CAtBat"     "CHits"      "CERuns"
## [11] "CRuns"      "CRBI"       "CWalks"     "League"     "DRA"
## [16] "PutOuts"    "Assists"    "Errors"     "Salary"     "Name"
```

```
Hitters = Hitters[which(!is.na(Hitters$Salary)), ]
```

Using the lasso to predict MLB salary: data prep

```
xfactors = model.matrix(Salary ~ League + Division +  
                        interaction(League, Division),  
                        data = Hitters)[-1]  
head(xfactors, 1)
```

```
##           LeagueN DivisionW interaction(League, Division)  
## -Alan Ashby           1           1  
##           interaction(League, Division)A.W  
## -Alan Ashby           0  
##           interaction(League, Division)N.W  
## -Alan Ashby           1
```


Using the lasso to predict MLB salary: fitting the lasso

```
x = as.matrix(data.frame(Hitters[, -c(14, 15, 19, 20)],  
                          xfactors))  
MLBsalary.lasso.cvfit = cv.glmnet(x, Hitters$Salary,  
                                   alpha = 1)  
MLBsalary.lasso.cvfit$lambda.min
```

```
## [1] 2.220313
```

Using the lasso to predict MLB salary: optimal λ for lasso

```
plot(MLBsalary.lasso.cvfit)
```



Using the lasso to predict MLB salary: lasso coefficients

```
head(coef(MLBsalary.lasso.cvfit), 12)
```

```
## 12 x 1 sparse Matrix of class "dgCMatrix"
##                               1
## (Intercept) 144.3797046
## AtBat      .
## Hits       1.3638038
## HmRun      .
## Runs       .
## RBI        .
## Walks      1.4973110
## Years      .
## CAtBat     .
## CHits      .
## CHmRun     .
## CRuns      0.1527517
```

Using the lasso to predict MLB salary: lasso coefficients

```
tail(coef(MLBsalary.lasso.cvfit), 10)
```

```
## 10 x 1 sparse Matrix of class "dgCMatrix"  
##                                1  
## CRBI                        0.32833941  
## CWalks                      .  
## PutOuts                     0.06625755  
## Assists                      .  
## Errors                      .  
## LeagueN                     .  
## DivisionW                   .  
## interaction.League..Division.N.E .  
## interaction.League..Division.A.W .  
## interaction.League..Division.N.W .
```



Using ridge regression to predict MLB salary: fitting

```
MLBsalary.rr.cvfit = cv.glmnet(x, Hitters$Salary,  
                                alpha = 0)  
MLBsalary.rr.cvfit$lambda.min
```

```
## [1] 28.01718
```

Using ridge regression to predict MLB salary: optimal λ

```
plot(MLBsalary.rr.cvfit)
```



Using ridge regression to predict MLB salary: coefficients

```
head(coef(MLBsalary.rr.cvfit), 12)
```

```
## 12 x 1 sparse Matrix of class "dgCMatrix"  
##              1  
## (Intercept) 2.051021e+02  
## AtBat       9.371095e-02  
## Hits        3.896188e-01  
## HmRun        1.222685e+00  
## Runs         6.227768e-01  
## RBI          6.184842e-01  
## Walks        8.085052e-01  
## Years        2.539892e+00  
## CAtBat       7.891144e-03  
## CHits        3.052107e-02  
## CHmRun       2.268930e-01  
## CRuns        6.116106e-02
```

Using ridge regression to predict MLB salary: coefficients

```
tail(coef(MLBsalary.rr.cvfit), 10)
```

```
## 10 x 1 sparse Matrix of class "dgCMatrix"
##                                     1
## CRBI                             0.063385593
## CWalks                           0.060508479
## PutOuts                           0.056023095
## Assists                           0.007589146
## Errors                           -0.173441672
## LeagueN                           2.334635145
## DivisionW                         -20.678801778
## interaction.League..Division.N.E  9.804002150
## interaction.League..Division.A.W -20.311680354
## interaction.League..Division.N.W -6.393999590
```

Making sense of ridge regression coefficients

```
sapply(c("AtBat", "Hits", "HmRun", "Runs", "RBI",  
        "Walks", "Years", "CAtBat", "CHits", "CHmRun",  
        "CRuns", "CRBI", "CWalks", "PutOuts", "Assists",  
        "Errors"),  
       function(x) mean(Hitters[,x])*  
       coef(MLBsalary.rr.cvfit)[x,])
```

##	AtBat	Hits	HmRun	Runs	RBI
##	37.8257316	42.0121640	14.2073188	34.0940699	31.8437034
##	Years	CAtBat	CHits	CHmRun	CRuns
##	18.5711494	20.9710590	22.0419025	15.7099711	22.0926294
##	CWalks	PutOuts	Assists	Errors	
##	15.7483095	16.2865315	0.9012905	-1.4904113	

Standard regression

First let's look at the full regression

```
x.full = data.frame(x, Salary = Hitters$Salary)
MLBsalary.lm = lm(Salary ~ ., data = x.full)
```

Standard regression

Table 1: Fitting linear model: $\text{Hitters\$Salary} \sim .$

	Estimate	Std. Error	t value	Pr(> t)
AtBat	-1.998	0.6353	-3.145	0.001866
Hits	7.551	2.381	3.172	0.001709
HmRun	4.412	6.208	0.7108	0.4779
Runs	-2.382	2.984	-0.7981	0.4256
RBI	-1.054	2.604	-0.4047	0.686
Walks	6.211	1.828	3.398	0.0007929
Years	-3.453	12.55	-0.2752	0.7834
CAtBat	-0.1689	0.1355	-1.246	0.2139
CHits	0.1093	0.6719	0.1626	0.8709
CHmRun	-0.2066	1.615	-0.128	0.8983
CRuns	1.478	0.747	1.978	0.04903
CRBI	0.8245	0.6917	1.192	0.2344
CWalks	-0.8133	0.3281	-2.479	0.01387
PutOuts	0.2819	0.07746	3.64	0.0003335
Assists	0.3721	0.2213	1.682	0.09395

Forward stepwise regression

```
MLBsalary.null = lm(Salary ~ 1, data = x.full)
MLBsalary.forward =
  stepAIC(MLBsalary.null, scope =
    list(upper = MLBsalary.lm,
         lower = MLBsalary.null),
    direction = "forward")
```

```
## Start:  AIC=3215.77
```

```
## Salary ~ 1
```

```
##
```

```
##
```

```
## + CRBI
```

```
## + CRuns
```

```
## + CHits
```

```
## + CAtBat
```

```
## + CHmRun
```

```
## + CWalks
```

```
## + RBI
```

	Df	Sum of Sq	RSS
--	----	-----------	-----

1	17139434	36179679
---	----------	----------

1	16881162	36437951
---	----------	----------

1	16065140	37253973
---	----------	----------

1	14759710	38559403
---	----------	----------

1	14692193	38626920
---	----------	----------

1	12792622	40526491
---	----------	----------

1	10771083	42548030
---	----------	----------

Forward stepwise regression

```
pander(MLBsalary.forward, style = "simple")
```

Table 2: Fitting linear model: $\text{Salary} \sim \text{CRBI} + \text{Hits} + \text{PutOuts} + \text{DivisionW} + \text{AtBat} + \text{Walks} + \text{CWalks} + \text{CRuns} + \text{CAtBat} + \text{Assists}$

	Estimate	Std. Error	t value	Pr(> t)
CRBI	0.7743	0.2096	3.694	0.0002706
Hits	6.918	1.647	4.201	3.686e-05
PutOuts	0.2974	0.07444	3.995	8.504e-05
DivisionW	-112.4	39.21	-2.866	0.004511
AtBat	-2.169	0.5363	-4.044	6.996e-05
Walks	5.773	1.585	3.643	0.0003274
CWalks	-0.8308	0.2636	-3.152	0.001818
CRuns	1.408	0.3904	3.607	0.0003731
CAtBat	-0.1301	0.0555	-2.344	0.01986
Assists	0.2832	0.1577	1.796	0.07367
(Intercept)	162.5	66.91	2.429	0.01583

Backward stepwise regression

```
MLBsalary.backward =  
  stepAIC(MLBsalary.lm, scope =  
    list(upper = MLBsalary.lm,  
         lower = MLBsalary.null),  
    direction = "backward")
```

```
## Start:  AIC=3046.13
```

```
## Salary ~ AtBat + Hits + HmRun + Runs + RBI + Walks + Year
```

```
##      CAtBat + CHits + CHmRun + CRuns + CRBI + CWalks + Pu
```

```
##      Assists + Errors + LeagueN + DivisionW + interaction
```

```
##      interaction.League..Division.A.W + interaction.Leagu
```

```
##
```

```
##
```

```
## Step:  AIC=3046.13
```

```
## Salary ~ AtBat + Hits + HmRun + Runs + RBI + Walks + Year
```

```
##      CAtBat + CHits + CHmRun + CRuns + CRBI + CWalks + Pu
```

```
##      Assists + Errors + LeagueN + DivisionW + interaction
```

```
##      interaction.League..Division.A.W
```

Backward stepwise regression

```
pander(MLBsalary.backward, style = "simple")
```

Table 3: Fitting linear model: $\text{Salary} \sim \text{AtBat} + \text{Hits} + \text{Walks} + \text{CAtBat} + \text{CRuns} + \text{CRBI} + \text{CWalks} + \text{PutOuts} + \text{Assists} + \text{DivisionW}$

	Estimate	Std. Error	t value	Pr(> t)
AtBat	-2.169	0.5363	-4.044	6.996e-05
Hits	6.918	1.647	4.201	3.686e-05
Walks	5.773	1.585	3.643	0.0003274
CAtBat	-0.1301	0.0555	-2.344	0.01986
CRuns	1.408	0.3904	3.607	0.0003731
CRBI	0.7743	0.2096	3.694	0.0002706
CWalks	-0.8308	0.2636	-3.152	0.001818
PutOuts	0.2974	0.07444	3.995	8.504e-05
Assists	0.2832	0.1577	1.796	0.07367
DivisionW	-112.4	39.21	-2.866	0.004511
(Intercept)	162.5	66.91	2.429	0.01583

Looks like the same model

```
length(names(MLBsalary.backward$coefficients))
```

```
## [1] 11
```

```
length(intersect(names(MLBsalary.backward$coefficients),  
                 names(MLBsalary.forward$coefficients)))
```

```
## [1] 11
```


Comparing forward and backward regression via LOOCV

```
forward_error
```

```
## -Alan Ashby
```

```
##      30060779
```

```
backward_error
```

```
## -Alan Ashby
```

```
##      28922347
```